

Performing Rolling Update Deployments



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

Understanding Rolling Update
Deployments

Creating a Rolling Update
Deployment

Rolling Update Deployment
in Action

Rolling Back Deployments

Rolling Back Deployments in Action



Kubernetes Resources



Storage/ConfigMaps/Secrets



Deployment

ReplicaSet



Pod



Container



Pod



Container



Pod



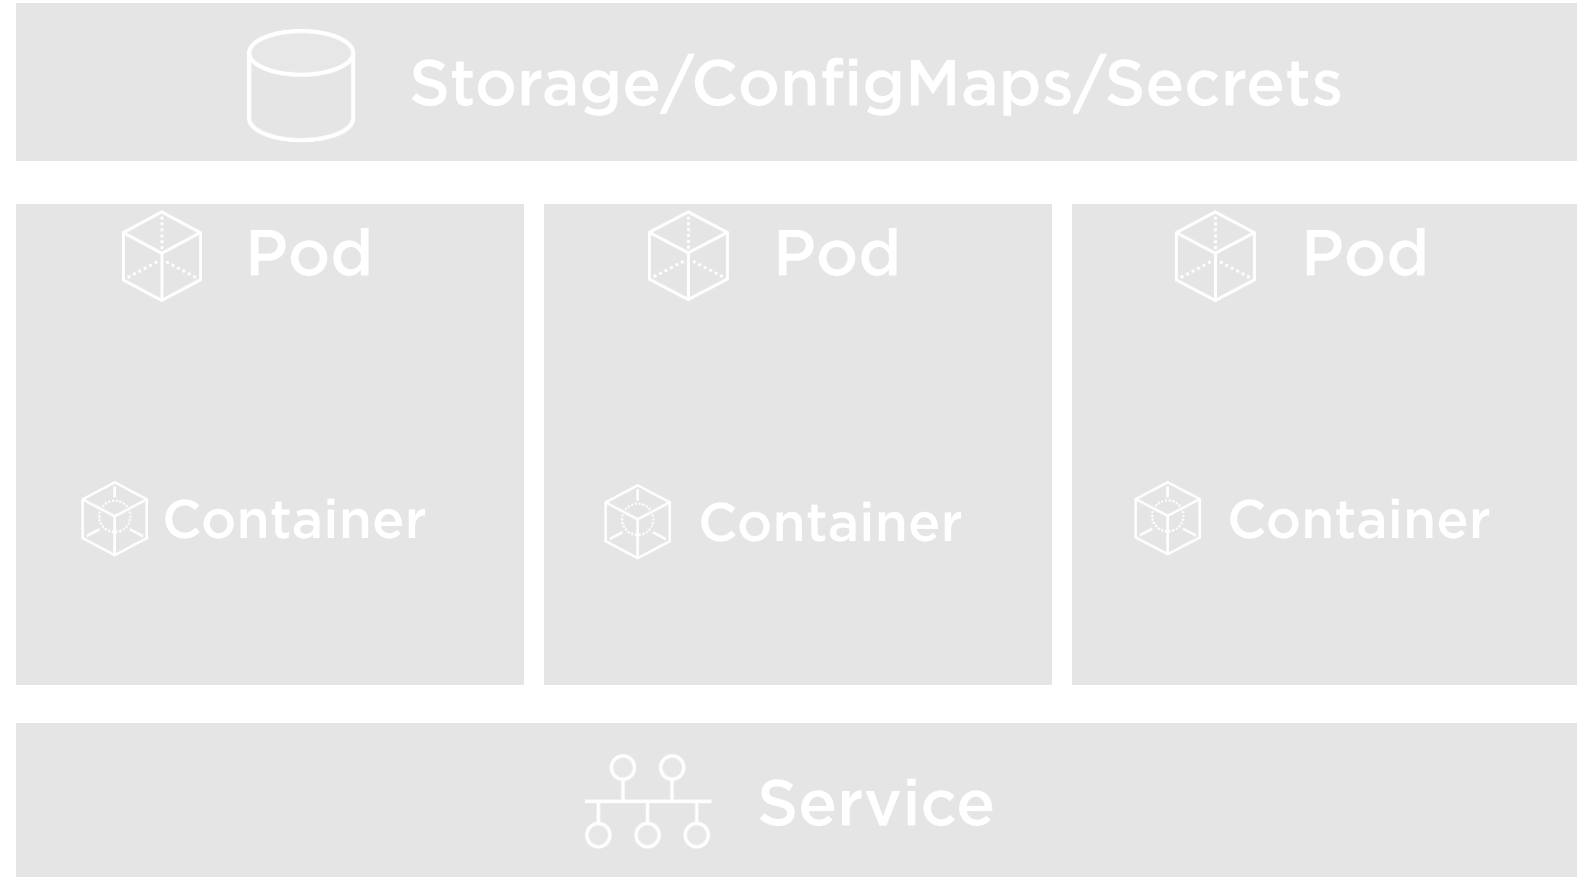
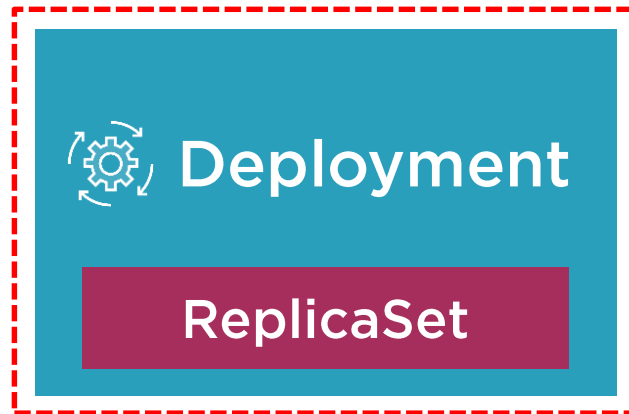
Container



Service



Kubernetes Resources



Kubernetes Resources



Storage/ConfigMaps/Secrets



Pod



Pod



Pod



Container



Container



Container



Service



Deployment

ReplicaSet



Understanding Rolling Update Deployments

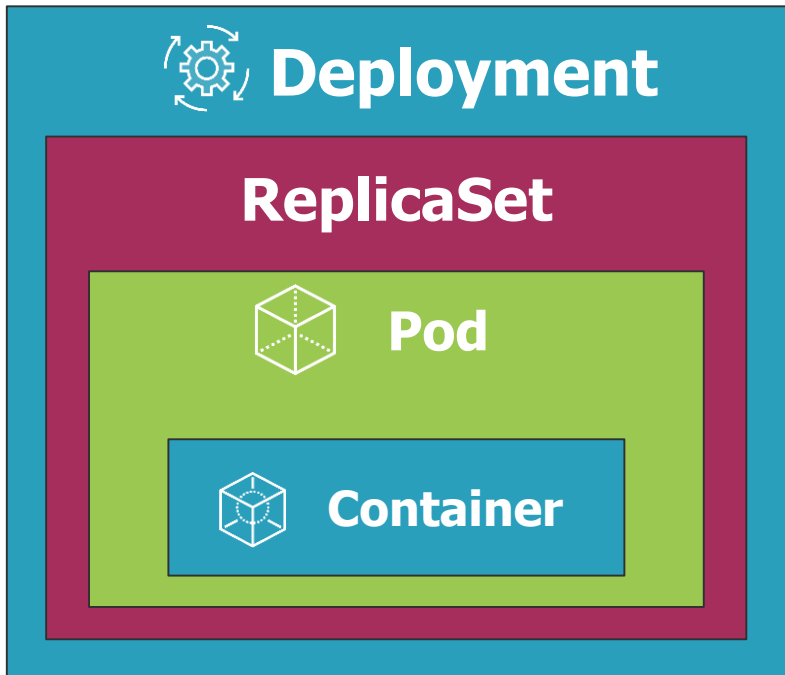


"Rolling updates allow Deployments' update to take place with zero downtime by incrementally updating Pods instances with new ones."

Kubernetes Documentation



Rolling Update Deployments



ReplicaSets increase new Pods while decreasing old Pods

Service handles load balancing traffic to available Pods

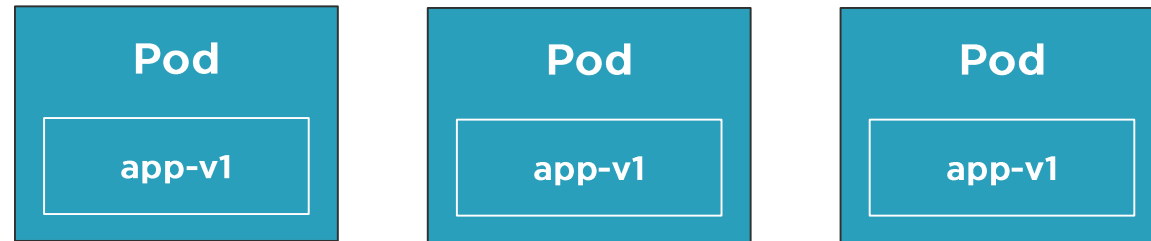
New Pods only scheduled on available Nodes

Deployments support two strategy options:

- Rolling Update (default and our focus here)
- Recreate (can result in down-time)

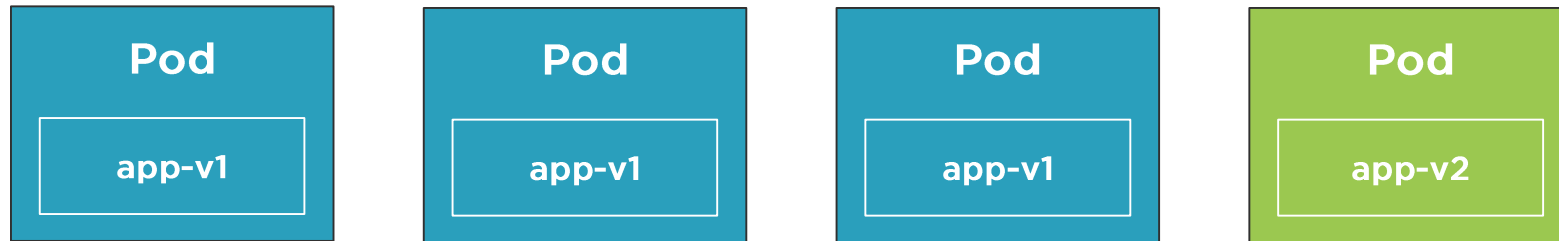
Rolling Update Deployments

Initial Pod State



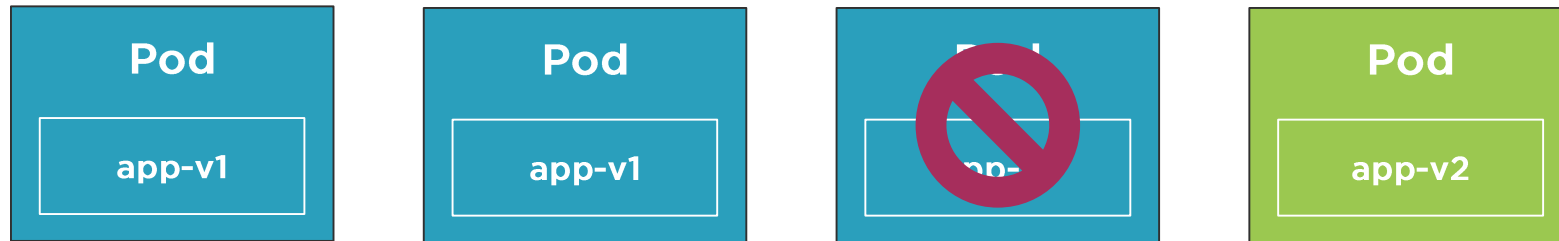
Rolling Update Deployments

Rollout New Pod



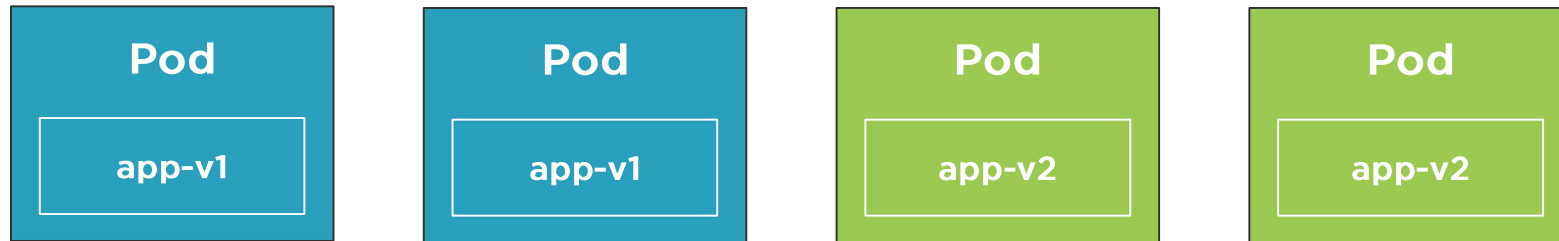
Rolling Update Deployments

Delete Pod



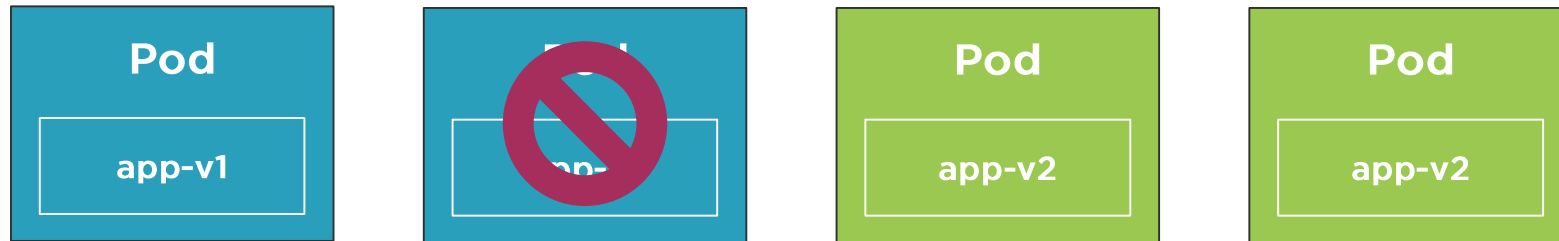
Rolling Update Deployments

Rollout New Pod



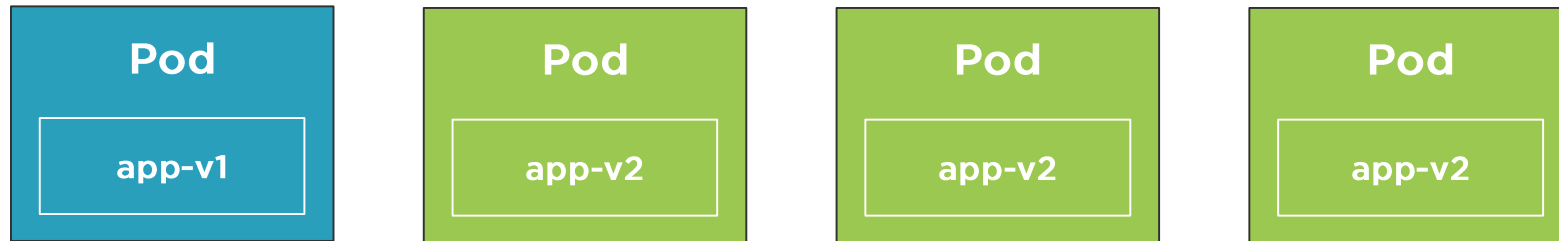
Rolling Update Deployments

Delete Pod



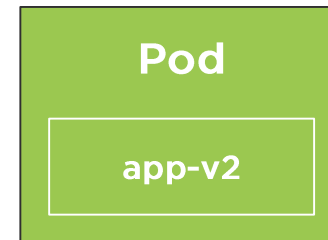
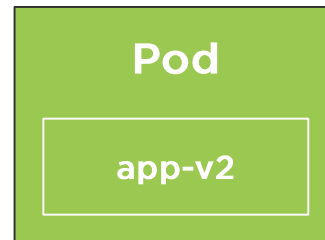
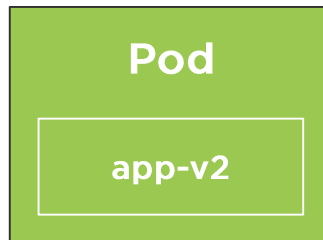
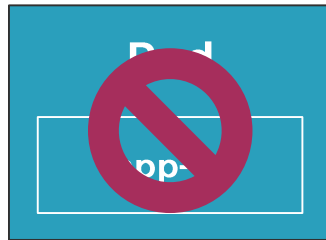
Rolling Update Deployments

Rollout New Pod



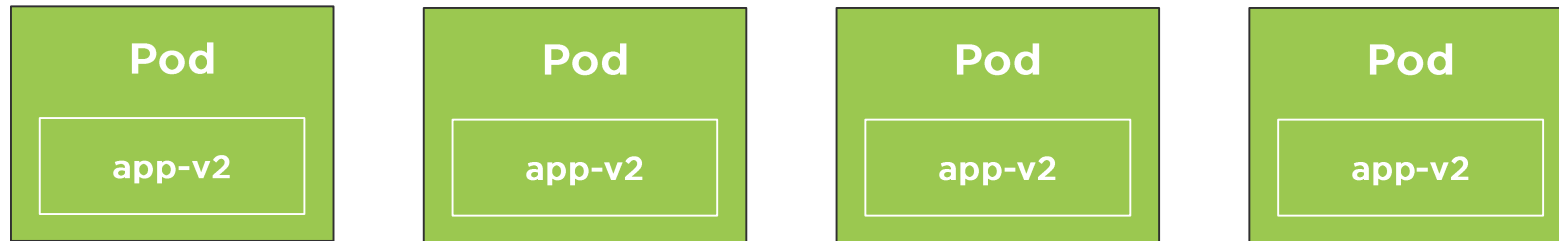
Rolling Update Deployments

Delete Pod

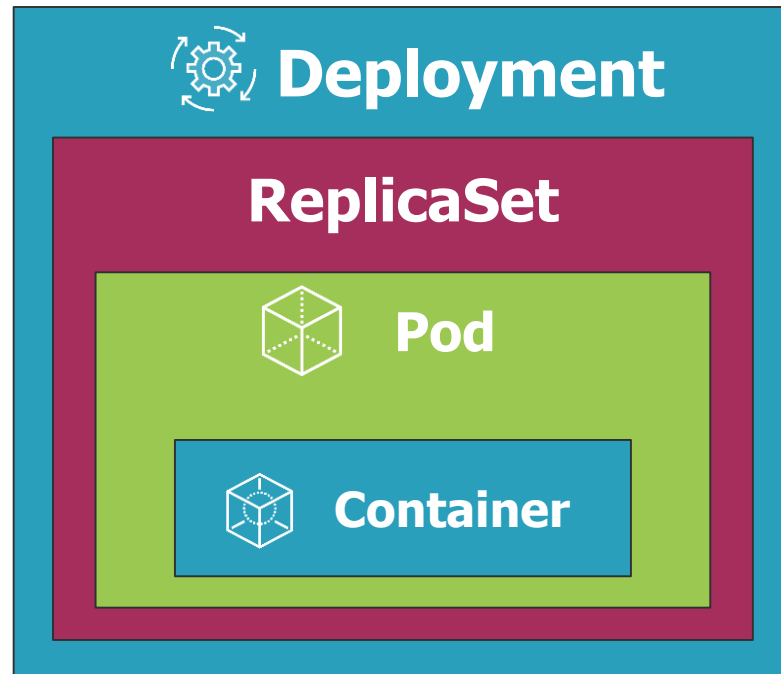


Rolling Update Deployments

Rollout New Pod



Deployments and Replicasets



Creating a Rolling Update Deployment



Defining a Rolling Update Deployment

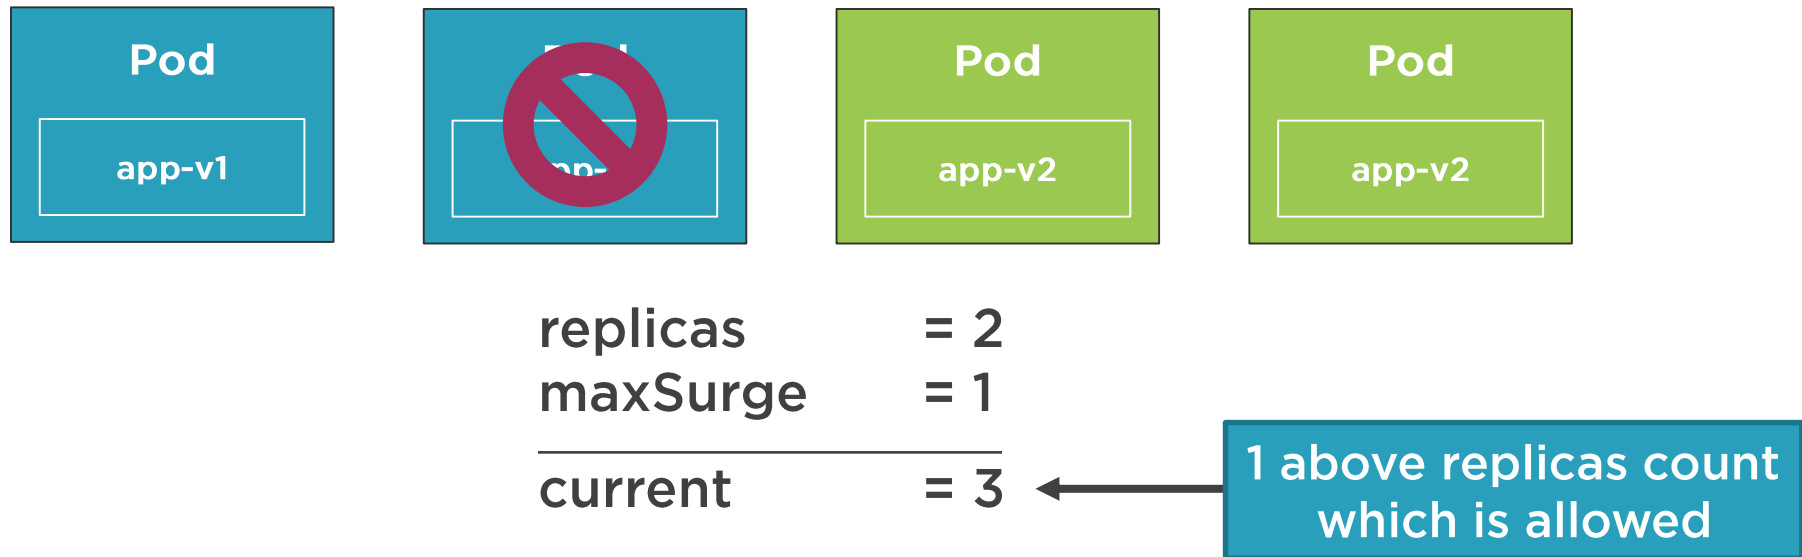
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
spec:
  replicas: 2
  minReadySeconds: 1
  progressDeadlineSeconds: 60
  revisionHistoryLimit: 5
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
  ...
```

- ◀ Number of Pod replicas
- ◀ Seconds new Pod should be ready to be considered healthy (0)
- ◀ Seconds to wait before reporting stalled Deployment
- ◀ Number of ReplicaSets that can be rolled back (10)
- ◀ RollingUpdate (default) or Recreate strategy
- ◀ Max Pods that can exceed the replicas count (25%)
- ◀ Max Pods that are not operational (25%)



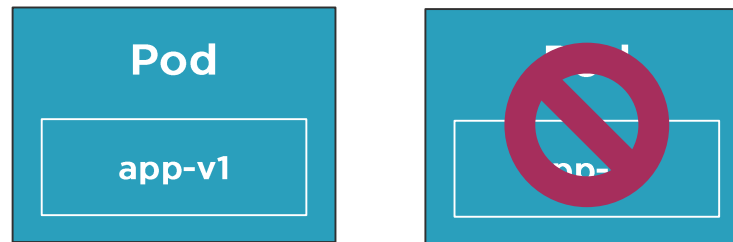
Understanding maxSurge

How many Pods can be added above the replicas count during the rolling update?



Understanding maxUnavailable

How many of the existing Pods can be made unavailable during a rolling update?



maxUnavailable = 1

It's OK for 1 of the 2
replicas to be unavailable



Creating the Deployment

Use the **kubectl create** command along with the **--filename** or **-f** switch

```
# Create initial deployment
```

```
kubectl create -f file.deployment.yml --save-config --record
```

Record the command
in the Deployment
revision history



Save configuration in
resource's annotations



Creating or Modifying a Deployment

Use the **kubectl apply** command along with the **--filename** or **-f** switch

Record the command
in the Deployment
revision history



```
# Create initial deployment
```

```
kubectl apply -f file.deployment.yml --record
```

Checking the Deployment Status

The **kubectl rollout status** command can be used to get information about a specific Deployment

```
# Get information about a Deployment
```

```
kubectl rollout status deployment [deployment-name]
```

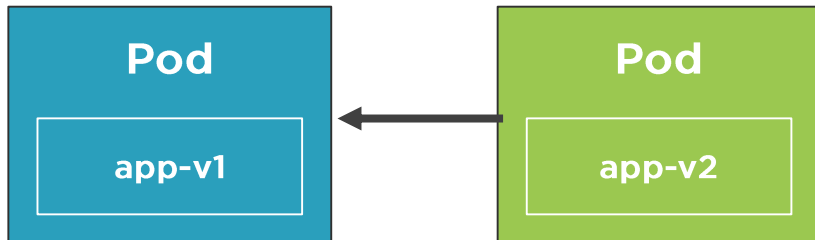

Rolling Update Deployment in Action



Rolling Back Deployments



Rolling Back Deployments



Rolling update revisions can be tracked using **--record**

If a Deployment has issues, a new Deployment can be applied, or you can revert to a previous revision

Several **kubectl** commands can be used for rollbacks:

- **kubectl rollout status**
- **kubectl rollout history**
- **kubectl rollout undo**

Checking Deployment History

The **kubectl rollout history** command can be used to view history of a Deployment

```
# Get information about a Deployment
```

```
kubectl rollout history deployment [deployment-name]
```

```
# Get information about a Deployment
```

```
kubectl rollout history deployment [deployment-name] --revision=2
```

Rolling Back a Deployment

Use the **kubectl rollout undo** command to rollback to a specific Deployment revision

```
# Check status
```

```
kubectl rollout status -f file.deployment.yml
```

```
# Rollback a Deployment
```

```
kubectl rollout undo -f file.deployment.yml
```

```
# Rollback to a specific revision
```

```
kubectl rollout undo deployment [deployment-name] --to-revision=2
```

Rolling Back Deployments in Action



Summary



Rolling updates are the default
Deployment strategy used by Kubernetes

Ensures zero-downtime during a
Deployment

Maximum and minimum Pods available
during a Deployment can be defined

Deployments can be recorded and stored
in history using `--record`

Deployments can be rolled back to a
specific revision

